

Counting 4-Cycles in Fully Dynamic Graphs: A Cyclic-Join Perspective

Anonymous Author(s)

Abstract

Counting subgraphs in a fully dynamic graph is used in databases for incremental view maintenance (IVM) of cyclic joins. The worst-case update time is settled for two of the three patterns on at most four vertices: triangles at $O(\sqrt{m})$, 4-cliques at $O(m)$. For 4-cycles the previous upper bound was $O(m^{2/3})$ [6]; Assadi and Shah [3] recently improved this to $O(m^{2/3-\epsilon})$ for a constant $\epsilon > 0$ using fast matrix multiplication (FMM), a tool previously thought inapplicable to dynamic IVM. This essay re-presents their algorithm through the IVM-for-cyclic-joins lens and contributes (i) a closed-form algebraic certificate that the underlying constraint system is feasible at the conjectural $\omega = 2$ with explicit rational parameters, complementing the numerical verification of [3]; and (ii) an empirical study contrasting the warm-up algorithm of [3, §3] (with NumPy GEMM substituting for FMM) against a layer-aware Simple-Wedge baseline, which locates the input regime — bilateral hubs on both sides of the cyclic ordering — where the class-routing machinery actually pays off.

CCS Concepts

• **Theory of computation** → **Dynamic graph algorithms**; *Streaming, sublinear and near linear time algorithms*; • **Information systems** → Database query processing.

Keywords

fully dynamic algorithms, subgraph counting, 4-cycles, fast matrix multiplication, incremental view maintenance, cyclic joins

1 Introduction

Counting occurrences of a small pattern graph (triangles, 4-cycles, 4-cliques) inside a larger host graph is one of the oldest and most broadly applicable problems in algorithmic graph theory. The same primitive underpins clustering and community-detection heuristics in social-network analysis [14], motif-discovery pipelines in computational biology [10], and recurring-subexpression evaluation in database query processing [13].

A further motivation, central to this paper, comes from database theory: *incremental view maintenance (IVM) for cyclic joins*. A cyclic join is a query $R_1 \bowtie R_2 \bowtie \dots \bowtie R_k$ over k binary relations whose attribute schemas form a cycle, with each relation sharing one attribute with its predecessor and one with its successor, and the last relation closing back to the first. Such queries are a classical hard case for join evaluation [11, 12]. The four-way instance considered throughout this paper is

$$J = A(L_1, L_2) \bowtie B(L_2, L_3) \bowtie C(L_3, L_4) \bowtie D(L_4, L_1),$$

which has a clean graph-theoretic reading: each binary relation is a bipartite edge set between two attribute domains, and $|J|$ equals the number of 4-cycles in the resulting 4-layered graph (Section 4.1). IVM updates a query’s answer incrementally after each

tuple change; maintaining $|J|$ under updates is therefore equivalent to dynamic 4-cycle counting (Section 4).

In each of these settings the host graph evolves under edge insertions and deletions, and recomputing the count from scratch is infeasible at scale. This motivates the *fully dynamic* model of subgraph counting, in which the exact count must be reported after every update; we measure cost by the *worst-case update time*, the maximum work on any single update, which is strictly stronger than amortised time.

The complexity of dynamic subgraph counting on small patterns is well understood for two of the three patterns on at most four vertices, summarised in Table 1. Triangles have worst-case update time $\Theta(\sqrt{m})$ [7, 9] (the lower bound is conditional on the OMv conjecture); 4-cliques have complexity $\Theta(m)$ for combinatorial algorithms under the 4-clique conjecture [6]. The 4-cycle lies between these two extremes: Hanauer, Henzinger, and Hua [6] obtained $O(m^{2/3})$ by generalising the triangle bound’s wedge-maintenance idea, with the only matching lower bound the $\Omega(m^{1/2-\gamma})$ inherited from triangles.

Assadi and Shah [3] prove that the $O(m^{2/3})$ upper bound is not tight. They give an algorithm with worst-case update time $O(m^{2/3-\epsilon})$ for a constant $\epsilon > 0$ depending on the matrix-multiplication exponent ω , the smallest constant for which two $n \times n$ matrices can be multiplied in $O(n^\omega)$ time. The current best is $\omega < 2.371339$ [1], yielding $\epsilon \approx 0.0098$; under the conjectural optimum $\omega = 2$ the bound improves to $\epsilon = 1/24$. The result is conceptually surprising: fast matrix multiplication (FMM) was previously believed incompatible with dynamic IVM (Section 2.4), and the Online Matrix-Vector multiplication (OMv) lower bound of [7] rules out only combinatorial improvements, leaving room for an algebraic speedup [3, §1].

Pattern	Upper bound	Lower bound
3-cycle (triangle)	$O(\sqrt{m})$ [9]	$\Omega(m^{1/2-\gamma})$ [7]
4-cycle	$O(m^{2/3-\epsilon})$ [3]	$\Omega(m^{1/2-\gamma})$ [7]
4-clique	$O(m)$ folklore [6]	$\Omega(m^{1-\gamma})$ [6]

Table 1: State of the art for fully dynamic subgraph counting on patterns of at most four vertices. Lower bounds are conditional (OMv for 3- and 4-cycles; combinatorial 4-clique conjecture for 4-cliques).

2 Related work

Dynamic subgraph counting builds on three lines of work: static FMM-based subgraph algorithms, dynamic wedge-maintenance for triangles, and the 4-vertex result of Hanauer, Henzinger, and Hua. The following sections present each of these in turn. We then discuss why fast matrix multiplication was long thought to be

incompatible with incremental view maintenance, and close with the positioning of the Assadi-Shah result.

2.1 Static algorithms for 4-cycles

The trivial static algorithm enumerates all *wedges* (length-2 paths $u-x-v$, i.e., pairs of edges sharing a common midpoint x) and, for each ordered pair (u, v) , counts the wedges between them; this gives $O(m^{3/2})$ time on graphs with m edges and is tight for combinatorial algorithms. Alon, Yuster, and Zwick [2] gave the first improvement to use fast matrix multiplication: their algorithm runs in time $O(m^{2\omega/(\omega+1)})$, which evaluates to roughly $O(m^{1.41})$ at the current ω and to $O(m^{4/3})$ at the conjectural optimum $\omega = 2$. The algorithm computes A^2 for the adjacency matrix A on the high-degree subgraph and combines it with a list-and-count step on the low-degree side; it was the first piece of evidence that FMM yields nontrivial speedups for 4-cycle counting in any setting. Crucially, this evidence was static: until 2025, no FMM-based speedup over the combinatorial bound was known in the dynamic regime.

2.2 Dynamic triangle counting and the OMv barrier

The triangle case is the template for this line of work. Kara, Ngo, Nikolic, Olteanu, and Zhang [9] maintain the triangle count in worst-case update time $O(\sqrt{m})$ by storing, for each ordered pair of vertices (u, v) , the number of common neighbours $W[u, v] = |N(u) \cap N(v)|$. The number of triangles created by inserting an edge (u, v) is exactly $W[u, v]$; updating W after the insertion costs $O(\deg(u))$ for the row indexed by u and $O(\deg(v))$ for the row indexed by v , and a high/low-degree partition (with threshold $\sqrt{m}$) caps the relevant degrees and gives the $\sqrt{m}$ bound. The OMv conjecture of Henzinger, Krinninger, Nanongkai, and Saranurak [7] implies that $\Omega(m^{1/2-\gamma})$ is the right answer; this lower bound holds even against algorithms that invoke fast matrix multiplication, because the OMv reduction is combinatorially robust. The same lower bound is inherited by 4-cycle counting via a straightforward reduction; [8] additionally show that the bound $\Omega(m^{1/2-\gamma})$ already holds for 4-cycle counting on random graphs, ruling out average-case improvements through that channel. Earlier work of [4, 5] maintained 3- and 4-vertex subgraph counts in $O(h)$ amortised update time, where h is the h -index of the graph; this is incomparable with bounds expressed purely in terms of m and is not a direct predecessor of the worst-case $O(\sqrt{m}) / O(m^{2/3})$ line of work, but it is the closest relative on the dynamic side.

2.3 The $O(m^{2/3})$ algorithm of Hanauer, Henzinger, and Hua

Hanauer, Henzinger, and Hua [6] extended the high/low-degree wedge-maintenance idea from triangles to all 4-vertex patterns. For 4-cycles, they partition the vertex set V into a *low-degree* class $V_L = \{v : \deg(v) \leq m^{1/3}\}$ and a *high-degree* class $V_H = V \setminus V_L$. The handshaking bound $\sum_v \deg(v) = 2m$ caps $|V_H|$ at $\frac{2m}{m^{1/3}} = O(m^{2/3})$, which is the load-bearing inequality of the analysis.

The algorithm maintains two indexed counts:

- $W_{lo}[u, v]$, the number of wedges $u-x-v$ with low-degree midpoint $x \in V_L$, stored for every pair of endpoints (u, v) ;

- $P_{lo}[u, v]$, the number of 3-paths $u-x-y-v$ with two low-degree midpoints $x, y \in V_L$, stored for every pair of endpoints.

On an update incident to u , the affected entries of W_{lo} are those of the form $W_{lo}[w, v]$ for a low-degree $w \in N(u) \cap V_L$ and a neighbour v of w ; iterating the inner pair costs at most $|N(u)| \cdot \max_{w \in V_L} |N(w)| \leq n \cdot m^{1/3}$ in the trivial reading, but only the low-degree neighbours of w contribute new wedges, and there are at most $m^{1/3}$ of them, so the work reduces to $|N(u)| \cdot m^{1/3} \leq m^{1/3} \cdot m^{1/3} = m^{2/3}$ when $u \in V_L$ and to $|N(u)| \cdot m^{1/3}$ in general. A symmetric analysis on P_{lo} keeps it within the same $O(m^{2/3})$ budget. Wedges through high-degree midpoints are maintained only between pairs of high-degree endpoints; there are $|V_H|^2 = O(m^{4/3})$ such pairs, but the per-update changes are confined to the at most $|V_H| = O(m^{2/3})$ pairs that share an endpoint with the updated edge.

Queries case-split on whether each midpoint of the counted 3-path is low or high. The bottleneck case, with two high-degree midpoints and at least one low-degree endpoint u , requires iterating over $N(u)$, which has size at most $m^{1/3}$, and combining stored counts; the inner step iterates over high-degree vertices, of which there are $O(m^{2/3})$. Either iteration alone gives a query cost of $O(m^{2/3})$, matching the maintenance budget.

This decomposition suggests that $O(m^{2/3})$ should be tight: 4-cycles would sit exactly at the geometric mean of triangles and 4-cliques, and any further improvement would require a fundamentally different technique. Assadi-Shah [3] show that this reading is mistaken: algebraic FMM techniques can push the upper bound below $m^{2/3}$.

2.4 Why FMM was thought not to help in IVM

Static FMM-based subgraph counting [2] requires both factors of the product, typically two slices of the adjacency matrix, to be assembled before the multiplication is invoked. In the fully dynamic regime the matrices are themselves the moving objects: each edge update changes a single entry, and waiting until enough updates accumulate to make a batched FMM call worthwhile only gives an *amortised* bound, not a worst-case one. The prevailing intuition was therefore that any FMM-based speedup must be amortised, and could not deliver a worst-case improvement of the kind we discuss [3, §1].

Assadi and Shah [3] circumvent this intuition by introducing *phases*: contiguous blocks of $m^{1-\delta}$ consecutive updates, for a constant $\delta > 0$. During a phase, queries are answered by a slower routine that uses precomputed products from the *previous* phase together with on-the-fly computation over the in-progress phase; in parallel, the FMM products needed for the *current* phase are computed incrementally, with the per-update share of the precomputation budgeted at $O(m^{2/3-\epsilon})$ work. By the end of the phase, the precomputed products are ready, and the next phase reuses them as its previous-phase data. Because both the per-update share of precomputation and the slower routine respect an $O(m^{2/3-\epsilon})$ ceiling on every individual update, not merely on average, the bound is genuinely worst-case (Section 3.2) rather than amortised.

2.5 Overview of the Assadi-Shah algorithm

The full Assadi-Shah algorithm runs on a 4-layered graph, groups the vertices in each layer into many degree classes (high, medium, low in the outer layers; dense, sparse in the inner layers), and pre-computes one rectangular FMM product per pair of compatible classes. A query then composes a constant number of stored entries plus a small-degree iteration. The ε in the exponent is the maximum value for which the system of constraints (one per class pair, expressing that both setup cost and per-query cost stay within the budget) is feasible at the prevailing rectangular FMM exponents. The argument also requires two simplifying assumptions which the paper later discharges: that every layer contains exactly $m^{2/3}$ vertices (no *tiny* vertices), and that vertices do not migrate between degree classes during a phase.

We cover the conceptual core: the reduction from general graphs to 4-layered graphs (Section 4), the simple $O(n)$ baseline (Section 5), which also serves as the algorithm we benchmark, and the restricted variant with two frozen relations (Sections 6–8) whose rectangular FMM step we develop in full and whose constraint system we verify in closed form. The de-assumption machinery (tiny vertices and vertex-class transitions) is left out of scope.

3 Preliminaries

3.1 Graphs and subgraph notation

Let $G = (V, E)$ be a simple, undirected, unweighted graph with no self-loops or parallel edges. We write $n = |V|$ and $m = |E|$. For a vertex $v \in V$, $\deg(v)$ denotes its degree and $N(v) = \{u : \{u, v\} \in E\}$ its neighborhood.

A k -*path* is a sequence of $k + 1$ distinct vertices $(v_0, v_1, \dots, v_k)$ with $\{v_i, v_{i+1}\} \in E$ for all $0 \leq i < k$. A k -*cycle* is a sequence $(v_0, \dots, v_{k-1})$ of k distinct vertices with $\{v_i, v_{i+1 \bmod k}\} \in E$ for all $0 \leq i < k$. A *wedge* is a 2-path (u, x, v) , equivalently a pair of edges $\{u, x\}, \{x, v\} \in E$ sharing the common midpoint x with distinct endpoints $u \neq v$. We write $W[u, v]$ for the number of wedges between u and v .

A *4-layered graph* is a graph whose vertex set partitions as $V = L_1 \cup L_2 \cup L_3 \cup L_4$, with each L_i an independent set and edges only between consecutive layers L_i, L_{i+1} (and between L_4 and L_1 , closing the cycle). The four edge sets are encoded by biadjacency matrices $A \subseteq L_1 \times L_2$, $B \subseteq L_2 \times L_3$, $C \subseteq L_3 \times L_4$, and $D \subseteq L_4 \times L_1$. A *layered k -path* (resp. *layered k -cycle*) is a k -path (resp. k -cycle) with one vertex in each consecutive layer. We index layers modulo 4 throughout.

3.2 Dynamic graph model

The graph evolves under a sequence of updates $u_1, u_2, \dots$, where each u_i is either $\text{INSERT}(e)$ or $\text{DELETE}(e)$ for an edge e on a fixed vertex set. Let G_i be the graph after the i -th update, with G_0 the empty graph. The model is *fully dynamic* because both insertions and deletions are permitted. Immediately after each update, the algorithm must output the exact number of 4-cycles in G_i .

The complexity measure is the *worst-case update time*: the maximum time spent processing any single update, taken over every update index and every legal sequence. This is strictly stronger than *amortized* time, which bounds only the average over a sequence; the headline bound of [3] is worst-case.




Figure 1: A 4-layered graph encoding the cyclic join $A \bowtie B \bowtie C \bowtie D$. Each binary relation is a bipartite edge set between two attribute domains. The highlighted layered 4-cycle (a, b, c, d) corresponds to a tuple in J . Adapted from [3], Fig. 1.

3.3 Fast matrix multiplication

Let ω denote the matrix-multiplication exponent: two $n \times n$ matrices can be multiplied in $O(n^\omega)$ arithmetic operations. The current record is $\omega < 2.371339$ [1], improving on Strassen's $\omega < 3$ [15]; whether $\omega = 2$ remains open. For rectangular products, $\omega(a, b, c)$ denotes the exponent of multiplying an $n^a \times n^b$ matrix by an $n^b \times n^c$ matrix in $n^{\omega(a, b, c)}$ time; the inner dimension n^b is shared by the two factors, while n^a and n^c are the outer dimensions of the left and right factors respectively, so $\omega(1, 1, 1) = \omega$. Improved bounds across the asymmetric range are also in [1]. An algorithm is *combinatorial* if it does not exploit Strassen-style identities; combinatorial algorithms therefore do not benefit from $\omega < 3$.

The use of FMM in subgraph counting rests on the standard identification of biadjacency-matrix products with path counts: if X is the biadjacency matrix between layers L_i and L_{i+1} and Y is the biadjacency matrix between L_{i+1} and L_{i+2} , then $(XY)_{u, w}$ equals the number of 2-paths from $u \in L_i$ to $w \in L_{i+2}$ through L_{i+1} . Iterating the product, $(XYZ)_{u, t}$ counts the 3-paths from u to t through three consecutive layers, which is exactly the quantity that drives 4-cycle queries in the layered setting (Theorem 1). Restricting X, Y, Z to specific degree classes yields rectangular factors whose products can be computed with the rectangular FMM bound $\omega(a, b, c)$.

3.4 The OMv conjecture

The *Online Matrix-Vector multiplication (OMv) conjecture* of [7] concerns the following problem: an algorithm preprocesses an $n \times n$ Boolean matrix M in polynomial time, then receives n Boolean vectors $v_1, \dots, v_n$ one at a time and must output Mv_i before seeing v_{i+1} . The conjecture asserts that no algorithm solves the entire instance in total time $O(n^{3-\gamma})$ for any constant $\gamma > 0$. The $\Omega(m^{1/2-\gamma})$ lower bounds for dynamic triangle and 4-cycle counting (Section 2.2) are reductions from OMv; we cite the conjecture only to motivate the lower bound and do not invoke it in any proof.

4 The cyclic-join equivalence

4.1 Joins as layered graphs

The motivating object behind 4-cycle counting on a 4-layered graph is the cyclic four-way join

$$J = A(L_1, L_2) \bowtie B(L_2, L_3) \bowtie C(L_3, L_4) \bowtie D(L_4, L_1),$$

familiar in database theory as a fundamental hard case for join evaluation [11, 12]. The four binary relations A, B, C, D over the attribute domains $L_1, \dots, L_4$ are exactly the biadjacency matrices of Section 3.1: a tuple in the relation between L_i and L_{i+1} corresponds to an edge between the two consecutive layers in $G' = (L_1 \cup L_2 \cup L_3 \cup L_4, E')$. A satisfying assignment $(a, b, c, d) \in L_1 \times L_2 \times L_3 \times L_4$ of J is precisely a closed sequence $a \rightarrow b \rightarrow c \rightarrow d \rightarrow a$ that visits one vertex per layer, that is, a layered 4-cycle in G' . The size of the join therefore equals the number of layered 4-cycles in G' . Figure 1 illustrates the correspondence on a small instance. Maintaining $|J|$ under tuple updates is then exactly the incremental view-maintenance problem for this cyclic join, which is the database-theoretic reading of our counting task.

4.2 Lifting a general graph to a 4-layered graph

From a general graph $G = (V, E)$, build a 4-layered graph $G' = (V', E')$ by replicating V in each layer: $V' = V \times \{1, 2, 3, 4\}$, with the i -th copy of V playing the role of L_i . Every edge $\{u, v\} \in E$ is lifted into all four biadjacency matrices: each of A, B, C, D records the pair (u, v) (equivalently, each layer pair $L_i \rightarrow L_{(i \bmod 4) + 1}$ gains the two directed edges (u_i, v_{i+1}) and (v_i, u_{i+1}) that encode the undirected edge between adjacent-layer copies). The construction is depicted in Figure 2.

Crucially, a single update (u, v) in G becomes four updates in G' , one per matrix, and the order in which these are applied is critical for the proof of Theorem 1. We adopt the following *update protocol* [3]:

Insertion. Insert (u, v) first into D , then into C , then into B , and finally into A .

Deletion. Delete (u, v) from A first, then from B , then from C , and last from D .

Query. Read off the change in the layered 4-cycle count immediately after the D -update, between the D step and the next of A, B, C .

In both cases, the protocol guarantees that the lifted edge (u, v) is *absent* from A, B , and C at the moment of query: on an insertion the edge has not yet propagated past D , and on a deletion it has already been removed from A, B, C before D . This asymmetry between D and the other three matrices is exactly what forces every length-3 *walk* counted by the algorithm to be a length-3 *path* in the underlying graph; without it, a walk could revisit u or v via the freshly written edge and the count would over-report.

4.3 Equivalence theorem

Theorem 1 (General-to-layered equivalence). *Let $\{u, v\}$ be the edge currently being updated in G , and let G' be the lifted graph of Section 4.2 immediately after the D -step of the protocol (and before any subsequent A, B, C step). At this moment, the number of length-3 walks in G' from $(u, 1) \in L_1$ to $(v, 4) \in L_4$ equals the number of length-3 paths from u to v in the current G .*

PROOF SKETCH. We follow [3]. The forward direction is by construction: each length-3 path u, x, y, v in G lifts to the walk $(u, 1), (x, 2), (y, 3), (v, 4)$ in G' , using one edge of A, B, C each. For the reverse direction, fix such a walk and write its underlying sequence as u, x, y, v . Simplicity of G (no self-loops to lift) rules out $x \in \{u, y\}$ and $y = v$. The




Figure 2: Lifting a general graph G (left) to a 4-layered graph G' (right). Every vertex of V is replicated in each layer $L_1, \dots, L_4$; gray edges schematically depict the lifted copies between consecutive layers. The 4-cycle u, v, x, y in G corresponds to the highlighted layered 4-cycle $(u, 1) \rightarrow (v, 2) \rightarrow (x, 3) \rightarrow (y, 4) \rightarrow (u, 1)$, with the closing $L_4 \rightarrow L_1$ edge drawn as a curved arc. For legibility, only the highlighted cycle is shown in full; remaining lifted edges are partial.

protocol rules out $x = v$ and $y = u$: each would require the edge $\{u, v\}$ to be present in A or C at query time, but the D -first ordering keeps (u, v) absent from A, B, C throughout the query window. The four vertices are therefore distinct, and the bijection follows. $\square$

4.4 Consequence

Theorem 1 reduces the dynamic 4-cycle counting problem on a general graph to dynamically counting layered 4-cycles on a 4-layered graph: the number of 4-cycles through a new edge (u, v) is exactly the layered 3-walk count delivered by the lifted instance under the protocol of Section 4.2, and the lift incurs only a constant blow-up in the number of layered updates per general-graph update. The restricted algorithm of Section 6 therefore works directly on a 4-layered graph with biadjacency matrices A, B, C, D . The simple $O(n)$ baseline of Section 5 is presented in the original general-graph setting and uses the same query-then-update timing idea adapted to its single wedge matrix.

5 A simple $O(n)$ algorithm

5.1 The wedge-counting idea

The simple algorithm of this section operates directly on a general graph $G = (V, E)$ on n vertices, without invoking the layered reduction of Section 4. The starting observation is that an edge update only changes the 4-cycle count through cycles *containing the updated edge*: if the count before an insertion of (u, v) is C_{old} , the count afterwards is $C_{\text{old}} + \Delta$, where Δ is the number of 4-cycles in the new graph that use the edge (u, v) . Each such 4-cycle is the closure $u \rightarrow w \rightarrow x \rightarrow v \rightarrow u$ of a length-3 path from u to v in the graph $G \setminus \{(u, v)\}$ (with the edge removed), so

$$\Delta = \#\{\text{length-3 paths from } u \text{ to } v \text{ in } G\}.$$

A length-3 path decomposes as a first edge $u \rightarrow w$ with $w \in N(u)$ followed by a length-2 path (a *wedge*) from w to v . Let $W[w, v]$ denote the number of wedges between w and v in G . Then

$$\Delta = \sum_{w \in N(u)} W[w, v]. \quad (1)$$

Algorithm 1 Maintain 4-cycle count via wedge matrix W .

```

1: procedure INSERT( $u, v$ )
2:    $\Delta \leftarrow \sum_{w \in N(u)} W[w, v]$  ▷ query
3:    $C \leftarrow C + \Delta$ 
4:   for  $w \in N(u)$  do  $W[w, v] += 1$ ;  $W[v, w] += 1$ 
5:   end for
6:   for  $w \in N(v)$  do  $W[w, u] += 1$ ;  $W[u, w] += 1$ 
7:   end for
8:   add  $(u, v)$  to  $E$  ▷ after query, to keep  $W$  free of  $(u, v)$ 
   during query
9: end procedure
10: procedure DELETE( $u, v$ )
11:   remove  $(u, v)$  from  $E$ 
12:   for  $w \in N(u)$  do  $W[w, v] -= 1$ ;  $W[v, w] -= 1$ 
13:   end for
14:   for  $w \in N(v)$  do  $W[w, u] -= 1$ ;  $W[u, w] -= 1$ 
15:   end for
16:    $\Delta \leftarrow \sum_{w \in N(u)} W[w, v]$  ▷ query
17:    $C \leftarrow C - \Delta$ 
18: end procedure

```

In principle the right-hand side counts length-3 walks u, w, x, v rather than paths; the order-of-operations rule of Section 5.2 guarantees that at query time the edge (u, v) is absent from W , which forces every counted walk to be a path with four distinct vertices and makes identity (1) hold without any correction term. The algorithm therefore reduces 4-cycle maintenance to maintaining the wedge-count matrix W .

5.2 The algorithm

Algorithm 1 maintains a global counter C for the current 4-cycle count alongside the wedge matrix W . Each update performs two distinct kinds of work: a *query* that uses identity (1) to compute the cycles through the updated edge, and a *maintenance* step that adjusts W to reflect the new edge set. Both steps touch only the rows and columns indexed by neighbours of u and v .

The order in which these two steps are performed is critical, and is the general-graph specialisation of the update protocol of Section 4.2: in both settings the new edge must be *absent* from the relevant data structure at the moment of query. For correctness, W must not contain any wedge that uses the edge (u, v) at the moment the query is evaluated; otherwise the expression $\sum_{w \in N(u)} W[w, v]$ would count walks that revisit u or v , breaking identity (1). On an INSERT, the edge (u, v) is not yet in E when the procedure starts, so the query is evaluated *first*, and only afterwards is W updated to record the wedges newly created through the inserted edge. On a DELETE, the edge (u, v) is still in E at the start, so W is updated *first*, removing the wedges that used (u, v) , and the query is evaluated only once those wedges are gone. In both cases, W at query time encodes only wedges of the graph $G \setminus \{(u, v)\}$, which is what identity (1) requires.

5.3 Correctness and complexity

Lemma 2 (Correctness and complexity of Algorithm 1). *Algorithm 1 maintains the exact 4-cycle count of any fully dynamic graph*

in worst-case update time $O(n)$, where $n = |V|$ is the (fixed) vertex count.

PROOF SKETCH. *Complexity.* Each update modifies only the rows of W indexed by neighbours of u and v , touching at most $|N(u)| + |N(v)| \leq 2n$ entries; the query is an inner product of length $|N(u)|$. Total: $O(n)$ per update.

Correctness. The query computes the number of walks u, w, x, v recorded in W at query time. The five potential vertex collisions $w \in \{u, v\}$ and $x \in \{u, v, w\}$ are ruled out as follows: simplicity of G excludes $w = u$, $x = w$, and $x = v$; the timing rule of Section 5.2 keeps $(u, v) \notin E$ at query time, which excludes $w = v$ and $x = u$ (each would require the absent edge to be present in W). Every counted walk is therefore a length-3 path from u to v , and conversely every such path is counted exactly once by the unique $w \in N(u)$ that witnesses its first hop. See [3, App. A] for the full case analysis. $\square$

5.4 Limitations of the $O(n)$ bound

The $O(n)$ bound of Lemma 2 is tight for sparse graphs in the worst case: when an update touches a vertex of degree $\Theta(n)$, every entry of the corresponding row of W may need to be inspected. For sparse graphs with $m = \Theta(n)$ this matches the natural target of expressing the update time purely as a function of m , since $n = \Theta(m)$. For denser graphs the bound is loose: when $n^2 = \omega(m)$, an update time of $O(n)$ can be much larger than any reasonable function of m , and the algorithm performs work that is not justified by the size of the current graph. A second, practical, drawback is the memory footprint: the wedge matrix W has n^2 entries, ≈ 16 MB at $n = 2000$ with 32-bit counters but ≈ 4 GB at $n = 32,000$ regardless of edge density, becoming prohibitive at the scale of modern social and web graphs.

These two shortcomings motivate the search for an algorithm whose update time is bounded purely as a function of the current edge count m . The first such bound was the $O(m^{2/3})$ algorithm of [6], which already uses a wedge-based approach as its starting point but partitions vertices by degree to break the $O(n)$ barrier. Section 6 develops the restricted algorithm of [3], whose update time is $O(m^{2/3-\epsilon_1})$ under simplifying assumptions and which is the building block for the $O(m^{2/3-\epsilon})$ result that finally crosses the $m^{2/3}$ threshold.

6 A restricted variant: fixed A and C

This section presents the restricted variant of the Assadi-Shah algorithm: a stripped-down version that already contains the central trick (precomputed FMM products combined with class-restricted iteration) but in which two of the four biadjacency matrices are frozen and two simplifying assumptions are imposed. The de-assumption machinery, occupying §6 and §7 of [3], is left out of scope; we sketch the move from the restricted variant to the full algorithm in Section 8.

6.1 Setup and assumptions

The restricted variant of [3] solves a simplified form of the layered counting problem in which two of the four biadjacency matrices are frozen: edge updates affect only B and D , while A and C are

fixed in advance. By symmetry it suffices to handle the query case in which the latest update is to D and we count new 4-cycles by counting 3-paths from $u \in L_1$ to $v \in L_4$ through A, B, C .

Two further structural assumptions support the analysis. First, the number of vertices in each layer satisfies $n \leq m^{2/3+2\epsilon}$, where $\epsilon > 0$ is a small constant fixed below; this rules out pathological layers that are much larger than the edge set warrants. Second, vertices do not migrate between degree classes during the algorithm; the bookkeeping needed to lift this assumption is the subject of a separate part of the source paper [3, §6–7] and is out of scope here. Under these assumptions the goal is a worst-case update time of $O(m^{2/3-\epsilon_1})$ for some constant $\epsilon_1 > 0$, matching the running time we will need when this algorithm is later invoked as a subroutine in the full result.

6.2 Vertex degree classes

The algorithm partitions vertices in the outer layers L_1 and L_4 into three degree classes, measured against the (fixed) matrices A and C respectively. A vertex is *High* (H) if its degree lies in $[m^{2/3-\epsilon_1}, n]$, *Medium* (M) if it lies in $[m^{1/3+\epsilon_1}, 2m^{2/3-\epsilon_1}]$, and *Low* (L) otherwise; the slight overlap between the M and H ranges keeps class boundaries on a single threshold rather than at an exact cutoff. We use A^{H*}, A^{M*}, A^{L*} to denote the row-restrictions of A to each class, and symmetrically C^{*H}, C^{*M}, C^{*L} for column-restrictions of C .

Inner-layer vertices in L_2 and L_3 are classified per chunk (§6.3): inside chunk B_i , a vertex is *Sparse* (S) if its degree in B_i is at most $m^{1/3-\epsilon_2}$ and *Dense* (D) otherwise. We correspondingly write $B_{i,DD}, B_{i,SS}, B_{i,DS}, B_{i,SD}$ for the four submatrices of B_i obtained by restricting both endpoints by class. The previous $O(m^{2/3})$ algorithm [6] used only two outer classes (high and low); the third, medium, class is what makes room for a rectangular FMM product in the restricted algorithm [3], since it lets the dense case ($A^{L*} \cdot B_{i,DD}$) sit on dimensions that current rectangular bounds can exploit.

6.3 Chunks and lazy maintenance

The stream of edge updates to B is grouped into *chunks* $B_1, B_2, \dots$ of size $m^{2/3-\epsilon_1}$, and we write $B_{<i} := \sum_{j<i} B_j$ for the aggregate of all completed chunks. Cheap data structures are maintained eagerly: each individual update touches only $O(m^{1/3+\epsilon_1})$ rows or columns, well within budget. Expensive data structures, the matrix products that require fast matrix multiplication, are not affordable per-update; we instead compute them once per chunk and amortize. Concretely, the products needed for chunk B_i are computed during the $m^{2/3-\epsilon_1}$ updates of the next chunk B_{i+1} , using a budget of $O(m^{4/3-2\epsilon_1})$ total work and therefore $O(m^{2/3-\epsilon_1})$ per update.

While inside chunk B_{i+1} , queries are answered by combining the freshly available data structures for $B_{<i+1}$ with a *lazy evaluation* for the in-progress chunk: edges of B_i and B_{i+1} are scanned one by one against the query endpoints, contributing $O(1)$ time per edge. A subtle issue arises when an edge inserted in chunk B_j is deleted in a later chunk $B_{j'}$; following [3] this is handled by recording the deletion as a *negative* edge in $B_{j'}$, so that the arithmetic in $B_{<i}$ remains correct without retroactively editing earlier chunks.

Class	Stored products
High in L_1, L_4	$A^{H*} \cdot B_{<i}; A^{H*} \cdot B_{<i} \cdot C^{*H}; B_{<i} \cdot C^{*H}$
Medium in L_1, L_4	$A^{M*} \cdot B_{<i}; B_{<i} \cdot C^{*M}$
Low in L_1, L_4	$A^{L*} \cdot B_{<i,DD}, A^{L*} \cdot B_{<i,SS}, A^{L*} \cdot B_{<i,SD},$ and three symmetric products with C^{*L}

Table 2: Data structures maintained by the restricted algorithm. Adapted from [3], Table 1.

6.4 Data structures

Table 2 lists the matrix products maintained across the three outer-layer classes, all evaluated against the chunk-aggregate $B_{<i}$. Each such product encodes the count of two- or three-hop paths between the indicated vertex classes. For example, $A^{H*} \cdot B_{<i}$ records, for every High vertex $u \in L_1$ and every $z \in L_3$, the number of L_1 - L_2 - L_3 wedges from u to z that use only edges of completed chunks. In database terminology this is a select-project-aggregate over the join $A \bowtie B$ restricted to L_1^H . The triple product $A^{H*} \cdot B_{<i} \cdot C^{*H}$ similarly stores the number of 3-paths between any two High endpoints, computed as a single rectangular matrix product per chunk transition. Medium classes require only the one-sided products $A^{M*} \cdot B_{<i}$ and $B_{<i} \cdot C^{*M}$ (the third hop is iterated at query time), while Low classes need finer S/D decompositions of $B_{<i}$ so that the dense-block product $A^{L*} \cdot B_{<i,DD}$ can be computed once per chunk via FMM.

Lemma 3 (Update-time accounting for Table 2). *The worst-case time to update all data structures in Table 2 is $O(m^{2/3-\epsilon_1})$ per update, with $O(m^{4/3-2\epsilon_1})$ amortized over chunk transitions.*

6.5 Answering queries

With the data structures of Table 2 in place, a query (u, v) with $u \in L_1, v \in L_4$ is answered by case analysis on the classes of u and v [3, Lemma 3.8].

Both High. The precomputed triple product $A^{H*} \cdot B_{<i} \cdot C^{*H}$ stores the 3-path count between every pair of High endpoints, so the answer is read off in $O(1)$.

Mixed (HM, ML, HL, MM). We iterate over the at most $2m^{2/3-\epsilon_1}$ neighbours of the lower-class endpoint and look up two-hop counts in $A^{H*} \cdot B_{<i}$ or $A^{M*} \cdot B_{<i}$, paying $O(m^{2/3-\epsilon_1})$ per query.

Both Low. We iterate over the $2m^{1/3+\epsilon_1}$ neighbours of v and combine DD, SS, SD contributions from the stored Low-class products, with the DS contribution computed symmetrically on u 's side via $B_{i,DS} \cdot C^{*L}$. Cost: $O(m^{1/3+\epsilon_1})$.

Lazy evaluation handles edges of the active chunks B_i, B_{i+1} in an additional $O(m^{2/3-\epsilon_1})$ time. Combining these bounds with Lemma 3 yields the main statement of the restricted variant: under Assumptions 1–3, 4-cycles in a fully dynamic 4-layered graph can be counted in worst-case time $O(m^{2/3-\epsilon_1})$.

7 The FMM step and the constraint system

The restricted algorithm relies on a single nontrivial use of fast matrix multiplication: the dense-block product $A^{L*} \cdot B_{i,DD}$, evaluated once per chunk transition. Its cost is determined by a rectangular

FMM call whose dimensions are fixed by the size of L_1 and the dense subset of L_2 .

7.1 Cost of the rectangular FMM step

Theorem 4 (Cost of the rectangular FMM step [3, Claim 3.6]). *The matrix product $A^{L^*} B_{i,DD}$ can be computed in time $m^{\omega(2/3+2\varepsilon, 1/3-\varepsilon_1+\varepsilon_2, 1/3-\varepsilon_1+\varepsilon_2)}$.*

PROOF. We bound the effective dimensions of the two factors. The matrix A^{L^*} has at most $|L_1| \leq n \leq m^{2/3+2\varepsilon}$ rows by Assumption 1, and we restrict its columns to indices in L_2 that are dense within chunk B_i . By definition, a Dense vertex of L_2 has degree at least $m^{1/3-\varepsilon_2}$ in B_i , so the total number of dense vertices in L_2 is at most $|B_i|/m^{1/3-\varepsilon_2} = m^{(2/3-\varepsilon_1)-(1/3-\varepsilon_2)} = m^{1/3-\varepsilon_1+\varepsilon_2}$. Hence the relevant dimensions of A^{L^*} are $m^{2/3+2\varepsilon} \times m^{1/3-\varepsilon_1+\varepsilon_2}$. The factor $B_{i,DD}$ is the restriction of B_i to dense rows and dense columns, so by the same accounting it has dimensions $m^{1/3-\varepsilon_1+\varepsilon_2} \times m^{1/3-\varepsilon_1+\varepsilon_2}$. The product is therefore a rectangular matrix multiplication of shape $(2/3+2\varepsilon, 1/3-\varepsilon_1+\varepsilon_2, 1/3-\varepsilon_1+\varepsilon_2)$, which can be performed in time $m^{\omega(2/3+2\varepsilon, 1/3-\varepsilon_1+\varepsilon_2, 1/3-\varepsilon_1+\varepsilon_2)}$ by definition of the rectangular FMM exponent [1]. $\square$

7.2 The constraint system

Charging the cost of Theorem 4 against the per-chunk budget $m^{4/3-2\varepsilon_1}$ from Lemma 3 yields the inequality

$$\omega(2/3+2\varepsilon, 1/3-\varepsilon_1+\varepsilon_2, 1/3-\varepsilon_1+\varepsilon_2) \leq 4/3-2\varepsilon_1, \quad (2)$$

which is the constraint that forces ε_1 to be strictly positive. Together with the High-High analogue and three threshold orderings, the full system is

$$\omega(2/3+2\varepsilon, 1/3-\varepsilon_1+\varepsilon_2, 1/3-\varepsilon_1+\varepsilon_2) \leq 4/3-2\varepsilon_1,$$

$$\omega(1/3+\varepsilon_1, 2/3-\varepsilon_1, 1/3+\varepsilon_1) \leq 4/3-2\varepsilon_1,$$

$$3\varepsilon_1+2\varepsilon \leq \varepsilon_2, \quad \varepsilon_1 \leq 1/6, \quad \varepsilon_1-\varepsilon_2 \leq 1/3,$$

corresponding to Equations (5), (2), (6), (7), (8) of [3, §3.4]. With the current rectangular FMM bounds of [1] the system is feasible at $\varepsilon_1 \approx 0.0420$, $\varepsilon_2 \approx 0.1457$, and $\varepsilon \approx 0.0098$ (paper Appendix B). The next subsection gives a clean closed-form solution at the conjectural optimum $\omega = 2$.

7.3 Closed-form feasibility at $\omega = 2$

Lemma 5 (Closed-form feasibility at $\omega = 2$). *Under the hypothetical $\omega = 2$, the constraint system above admits the solution $\varepsilon_1 = 1/24$, $\varepsilon_2 = 5/24$, $\varepsilon = 1/24$.*

PROOF. At $\omega = 2$, square matrix multiplication runs in input-size time, and the rectangular exponent collapses to $\omega(a, b, c) = \max(a+b, b+c, a+c)$, the cost of merely reading the largest pair of factors. We verify each constraint in turn.

Threshold ordering, $\varepsilon_1 \leq 1/6$ (constraint (7) of [3]): $1/24 < 4/24 = 1/6$. *Threshold ordering*, $\varepsilon_1 - \varepsilon_2 \leq 1/3$ (constraint (8) of [3]): $1/24 - 5/24 = -4/24 < 1/3$. *Sparse-Sparse / Sparse-Dense iteration*, $3\varepsilon_1 + 2\varepsilon \leq \varepsilon_2$ (constraint (6) of [3]): $3 \cdot \frac{1}{24} + 2 \cdot \frac{1}{24} = \frac{5}{24} = \varepsilon_2$, so the constraint holds with equality. *Low-Dense FMM*, our Eq. (2) (constraint (5) of [3]): substitute $a = 2/3 + 2\varepsilon = 2/3 + 1/12 = 3/4$ and $b = c = 1/3 - \varepsilon_1 + \varepsilon_2 = 1/3 + 4/24 = 1/2$. Then

$$\omega(3/4, 1/2, 1/2) = \max(\frac{3}{4} + \frac{1}{2}, \frac{1}{2} + \frac{1}{2}, \frac{3}{4} + \frac{1}{2}) = \frac{5}{4}.$$

The right-hand side is $4/3 - 2\varepsilon_1 = 4/3 - 1/12 = 16/12 - 1/12 = 15/12 = 5/4$, so the constraint is satisfied with equality. *High-High product* (constraint (2) of [3]), $\omega(1/3+\varepsilon_1, 2/3-\varepsilon_1, 1/3+\varepsilon_1) \leq 4/3-2\varepsilon_1$: the dimensions evaluate to $a = c = 9/24 = 3/8$ and $b = 15/24 = 5/8$.

At $\omega = 2$,

$$\omega(3/8, 5/8, 3/8) = \max(\frac{3}{8} + \frac{5}{8}, \frac{5}{8} + \frac{3}{8}, \frac{3}{8} + \frac{3}{8}) = 1 \leq 5/4 = 4/3 - 2\varepsilon_1,$$

verified by the same method as our Eq. (2). All five constraints are satisfied, so the proposed assignment is feasible. $\square$

The verification highlights why FMM is essential. The two FMM constraints (constraints (2) and (5) of [3], the latter our Eq. (2)) bind tightly at $\omega = 2$, while the three threshold orderings (constraints (6), (7), (8)) have slack; without a sub-cubic exponent the FMM constraints would force $\varepsilon_1 = 0$. Appendix B of [3] carries out the analogous verification with the current numerical bounds of [1] and confirms a strictly positive slack.

8 Phases and the full algorithm

The restricted algorithm relies critically on the fact that A and C are frozen: the running aggregate $A^{L^*} \cdot B_{<i,DD} = A^{L^*} \cdot B_{<i-1,DD} + A^{L^*} \cdot B_{i-1,DD}$ that drives the chunk machinery only makes sense when the left factor does not change. Once A also receives updates, this telescoping breaks down and the FMM step has to be redesigned.

The full algorithm of [3] replaces chunks with *phases* of size $m^{1-\delta}$ for a constant $\delta > 0$. The mechanism splits each phase into three concurrent activities, all sharing the per-update budget of $m^{2/3-\varepsilon}$:

Slow path during the phase. Queries are answered using the precomputed products of the *previous* phase, combined with on-the-fly evaluation against the updates accumulated so far in the current phase. Each such query costs $O(m^{2/3-\varepsilon})$ in the worst case.

FMM precomputation in parallel. During the $m^{1-\delta}$ updates of the current phase, the algorithm amortises a single multiplication of two square matrices of dimension $m^{2/3+2\varepsilon}$ across the phase. The per-update share is $(m^{2/3+2\varepsilon})^\omega / m^{1-\delta}$ time, which the constraint below forces to be $O(m^{2/3-\varepsilon})$.

Roll-forward at phase boundaries. When the phase ends, the freshly computed products become the new “previous-phase” data structure, and a new phase begins.

Because the per-update cost of every concurrent activity is $O(m^{2/3-\varepsilon})$ on every update, not merely on average, the resulting bound is worst-case. The cost analysis in [3, §4] distills to the single inequality

$$1 - \delta \geq (2\omega + 1)\varepsilon + (\omega - 1)\frac{2}{3},$$

which has a feasible solution with $\delta > 0$ if and only if $\omega < 5/2 - O(\varepsilon)$. This is the conceptual content of the result: the speedup vanishes long before the cubic barrier, and Strassen’s $\omega \approx 2.81$ [15] is therefore *not* sufficient. Only the asymptotically aggressive bounds of [1], which give $\omega \approx 2.371339$, push the exponent below $5/2$ at all. At the conjectural $\omega = 2$ with $\varepsilon = 1/24$, the inequality binds at $\delta = 1/8$ (since $1 - 1/8 = 7/8 = 5/24 + 16/24 = (2 \cdot 2 + 1)/24 + 2/3$).

Theorem 6 (Main result of [3], restated). *There is an algorithm for counting 4-cycles in any fully dynamic graph in worst-case update*

time $O(m^{2/3-\epsilon})$ for some constant $\epsilon > 0$. With current $\omega = 2.371339$, $\epsilon \approx 0.0098$; with $\omega = 2$, $\epsilon = 1/24$.

We restate Theorem 6 verbatim from [3] and do *not* reprove it. Its proof combines the restricted algorithm of this section, the phase machinery sketched above, the removal of Assumption 1 via tiny-vertex handling [3, §6], and the removal of Assumption 2 via vertex-class transitions [3, §7], none of which we reproduce. The formal contribution of this paper is limited to Theorem 1 (the general-to-layered reduction), Lemma 2 (the $O(n)$ baseline), Theorem 4 (the rectangular FMM step), and Lemma 5 (feasibility of the constraint system at $\omega = 2$).

9 Experiments

The warm-up of Section 6 avoids $O(n)$ -per-update enumeration when an endpoint is high-degree by storing matrix products restricted to high-degree submatrices and dispatching by class. We ask empirically: when does this class-routing actually pay off in wall-clock time? Comparing WARMUP-V3 only to SIMPLE-WEDGE (which sees the union graph and pays $O(\deg_{\text{union}}(u))$ per update) is misleading: a layer-aware baseline already avoids much of that cost. We therefore introduce LAYERED-SIMPLE-WEDGE, the apples-to-apples baseline.

9.1 Setup

We compare three counters with a common insert/delete/count interface, all written in Python 3.11 with NumPy. SIMPLE-WEDGE is the algorithm of Section 5 as written: it maintains a dense union-graph wedge matrix W and pays $O(\deg(u) + \deg(v))$ per update. LAYERED-SIMPLE-WEDGE is the same idea, made aware of the four-layer structure: it maintains only the wedge counts $W_{13}[i, w] = \sum_j A[i, j]B[j, w]$ relevant to layered 4-cycles and answers a D -edge update by $\sum_{w \in N_C(l)} W_{13}[i, w]$. It counts exactly the same object as the warm-up ($|J| = \text{tr}(ABCD)$) and is the apples-to-apples baseline. WARMUP-V3 is a faithful-in-shape implementation of Section 6 with two simplifications, both stated in code/README.md: NumPy GEMM substitutes for fast rectangular matrix multiplication, and Assumptions 1+2 are fixed at construction so updates that would change a vertex’s class raise an error. Chunk size is fixed at one, collapsing the chunk-Dense class to empty.

Inputs are layered graphs built by `layered_hub`: a *planted-hub* family with $k_{\text{hubs}} = 5$ vertices in L_1 at degree linear in n , and a *bilateral-hub* family with 5 such hubs in both L_1 and L_4 — the only setting that triggers *HH*-class queries in WARMUP-V3. Both run a D -only stream biased toward hub vertices (75% for planted-hub, 85% for bilateral). Per-update times are medians over 200 updates across 3 random seeds with IQR error bars; build costs are reported separately, since WARMUP-V3 pays an $O(n^3)$ GEMM precompute that should not enter a per-update headline. A correctness gate (61 lockstep tests, all passing) verifies that $\text{WARMUP-V3.COUNT}() = \text{tr}(ABCD) = \text{LAYERED-SIMPLE-WEDGE.COUNT}()$ after every step of random B+D streams.

9.2 Two regimes

Hubs only on one side (Figure 3). At $n_{\text{per-layer}} = 2000$ with hubs only in L_1 : SIMPLE-WEDGE costs $\approx 420 \mu\text{s}/\text{update}$ (linear in n), WARMUP-V3 $\approx 4.1 \mu\text{s}/\text{update}$, LAYERED-SIMPLE-WEDGE $\approx 2.7 \mu\text{s}/\text{update}$




Figure 3: Planted-hub graphs ($k_{\text{hubs}} = 5$ in L_1 only). SIMPLE-WEDGE scales linearly in n (slope 0.93), the textbook $O(\deg_{\text{union}}(\text{hub}))$ worst case. The layered baseline beats WARMUP-V3 by $\approx 1.6\times$: with no hubs in L_4 , every D -update is *HL* or *LL* and the $O(1)$ *HH*-lookup never fires.




Figure 4: Bilateral-hub graphs ($k_{\text{hubs}} = 5$ in both L_1 and L_4). Most D -updates now hit the *HH* cell of WARMUP-V3’s dispatch table; the layered baseline still pays $O(\deg_C(L_4 \text{ hub})) \approx O(n)$. WARMUP-V3 wins, by a factor that grows with n .

— the layer-aware baseline beats the warm-up by $\approx 1.6\times$. The reason is structural: with hubs only in L_1 , every D -update query falls into the *HL* or *LL* class. The *HL* lookup in WARMUP-V3’s $A^{H*} \cdot B$ table iterates $|N_C(l)|$ entries — exactly the same row sum as the layered baseline, but with extra class-dispatch overhead in Python. WARMUP-V3’s distinguishing $O(1)$ *HH*-lookup is never triggered.

Hubs on both sides (Figure 4). Planting 5 hubs in L_4 as well pushes most D -updates into the *HH* cell of WARMUP-V3’s dispatch table, where the precomputed $A^{H*} \cdot B \cdot C^{*H}$ matrix yields an $O(1)$ lookup. At $n_{\text{per-layer}} = 2000$, LAYERED-SIMPLE-WEDGE costs $34.1 \mu\text{s}/\text{update}$ while WARMUP-V3 costs $6.5 \mu\text{s}/\text{update}$ — a $5.2\times$ speedup, growing with n . The layered baseline must iterate $|N_C(L_4 \text{ hub})| \Theta(n)$ entries; WARMUP-V3 reads one entry of its precomputed table. This is the mechanism promised by Section 6: the warm-up’s win

comes from precomputed matrix products at high-degree submatrices, and only when the input has hubs on both sides of the cyclic ordering.

9.3 Caveat

We implemented WARMUP-V3 faithfully *in shape* but with NumPy GEMM standing in for fast rectangular FMM (no practical implementation of the latter exists at our scale; even Strassen’s algorithm needs n in the thousands to break even with BLAS [1]). The $O(n^3)$ GEMM precompute is large in absolute terms: at $n_{\text{per-layer}} = 2000$ it takes 23 s, against 0.10 s for LAYERED-SIMPLE-WEDGE. Whether the warm-up’s per-update $O(m^{2/3-\epsilon_1})$ actually beats the layered baseline’s $O(m^{1/2})$ in wall-clock, with FMM in place, is the central practical question raised by [3]; our experiments cannot answer it because we do not have rectangular FMM. What our experiments *can* establish is that the warm-up’s class-routing mechanism, in isolation from its FMM ingredient, helps only on graphs whose hub structure makes *HH*-class queries fire frequently — which is precisely the structural assumption baked into the design of Section 6.

10 Conclusion and further work

We presented the recent algorithm of Assadi and Shah for dynamic 4-cycle counting through the lens of incremental view maintenance for the cyclic four-way join. We surveyed the line of work culminating in their result (Section 2), recalled the layered reduction that lets us study the problem on a 4-layered graph (Section 4), presented a simple $O(n)$ wedge-maintenance baseline (Section 5), and exposed the conceptual core of the Assadi-Shah algorithm: the vertex-class partition with its precomputed FMM products (Section 6), the rectangular FMM step that drives the parameter trade-off (Section 7), and the phase mechanism that turns those precomputed products into a worst-case guarantee (Section 8).

Our formal contribution is Lemma 5: a closed-form algebraic certificate that the constraint system at the heart of the restricted algorithm is feasible at the conjectural optimum $\omega = 2$, with explicit rational parameters $\epsilon_1 = 1/24$, $\epsilon_2 = 5/24$, and $\epsilon = 1/24$. Where the source paper verifies feasibility numerically with the current rectangular FMM bounds (Appendix B of [3]), the $\omega = 2$ closed form makes interpretable why two of the five constraints (the rectangular FMM constraints) bind tightly and the others have slack; this clarifies which inequalities are the load-bearing parts of the parameter trade-off. The supporting dimension analysis (Theorem 4) is taken from the source paper. In Section 9 we located, empirically, the input regime in which the warm-up’s class-routing actually pays off — graphs with hubs on both sides of the cyclic ordering — by contrasting it with a layer-aware baseline that already captures most of the warm-up’s benefit on simpler inputs.

We highlight three directions for further work. First, the gap between the $O(m^{2/3-\epsilon})$ upper bound of [3] and the $\Omega(m^{1/2-\gamma})$ OMv-conditional lower bound of [7]: where on $[1/2, 2/3]$ does the truth lie? The OMv reduction passes through boolean matrix-vector products and so does not rule out further FMM-based speedups, while [3] explicitly notes that a combinatorial $\Omega(m^{2/3})$ lower bound is not yet ruled out either, so the gap could close from either side. Second, IVM-with-FMM as a research direction: this is the first time fast matrix multiplication has been leveraged for join-size estimation

in dynamic graphs, and the authors suggest the technique could inspire a new class of IVM algorithms [3]. Where else does FMM break a “natural” barrier in dynamic IVM? Other cyclic joins, longer cycles, richer pattern queries? An adjacent direction is approximate counting under updates: [16] combines sampling with FMM for static approximate triangle counting, and a dynamic counterpart would close a parallel gap to the one studied here. Third, on the empirical side, a faithful implementation of the restricted algorithm with a competitive rectangular FMM routine would test whether the constants of [1] let the asymptotic improvement surface at realistic graph sizes; our benchmarks of the simple algorithm cannot speak to this, and we view it as the natural next step.

References

- [1] Josh Alman, Ran Duan, Virginia Vassilevska Williams, Yinzhao Xu, Zixuan Xu, and Renfei Zhou. 2025. More Asymmetry Yields Faster Matrix Multiplication. In *Proceedings of the 2025 Annual ACM-SIAM Symposium on Discrete Algorithms (SODA)*. 2005–2039.
- [2] Noga Alon, Raphael Yuster, and Uri Zwick. 1997. Finding and Counting Given Length Cycles. *Algorithmica* 17, 3 (1997), 209–223.
- [3] Sepehr Assadi and Vihan Shah. 2025. An Improved Fully Dynamic Algorithm for Counting 4-Cycles in General Graphs using Fast Matrix Multiplication. arXiv:2504.10748 [cs.DS] PODS 2025, Article 091. <https://doi.org/10.1145/3725228>.
- [4] David Eppstein, Michael T. Goodrich, Darren Strash, and Lowell Trott. 2012. Extended Dynamic Subgraph Statistics Using h-Index Parameterized Data Structures. *Theoretical Computer Science* 447 (2012), 44–52.
- [5] David Eppstein and Emma S. Spiro. 2009. The h-Index of a Graph and its Application to Dynamic Subgraph Statistics. In *Algorithms and Data Structures (WADS 2009) (LNCS, Vol. 5664)*. Springer, 278–289.
- [6] Kathrin Hanauer, Monika Henzinger, and Qi Cheng Hua. 2022. Fully Dynamic Four-Vertex Subgraph Counting. In *1st Symposium on Algorithmic Foundations of Dynamic Networks (SAND 2022) (LIPIcs, Vol. 221)*. Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 18:1–18:17.
- [7] Monika Henzinger, Sebastian Krimminger, Danupon Nanongkai, and Thatchaphol Saranurak. 2015. Unifying and Strengthening Hardness for Dynamic Problems via the Online Matrix-Vector Multiplication Conjecture. In *Proceedings of the Forty-Seventh Annual ACM Symposium on Theory of Computing (STOC)*. 21–30.
- [8] Xiao Hu, Qizhi Liu, and Stavros Sintos. 2022. Lower Bounds for Dynamic Maintenance of Cyclic Joins. In *Proceedings of the 41st ACM SIGMOD-SIGACT-SIGAI Symposium on Principles of Database Systems (PODS)*. 275–286.
- [9] Ahmet Kara, Hung Q. Ngo, Milos Nikolic, Dan Olteanu, and Haozhe Zhang. 2020. Maintaining Triangle Queries under Updates. *ACM Transactions on Database Systems* 45, 3 (2020), 11:1–11:46.
- [10] Ron Milo, Shai Shen-Orr, Shalev Itzkovitz, Nadav Kashtan, Dmitri Chklovskii, and Uri Alon. 2002. Network Motifs: Simple Building Blocks of Complex Networks. *Science* 298, 5594 (2002), 824–827.
- [11] Hung Q. Ngo, Ely Porat, Christopher Ré, and Atri Rudra. 2018. Worst-Case Optimal Join Algorithms. *J. ACM* 65, 3 (2018).
- [12] Hung Q. Ngo, Christopher Ré, and Atri Rudra. 2014. Skew Strikes Back: New Developments in the Theory of Join Algorithms. *ACM SIGMOD Record* 42, 4 (2014), 5–16.
- [13] Pedro Ribeiro, Pedro Paredes, Miguel E. P. Silva, David Aparicio, and Fernando Silva. 2021. A Survey on Subgraph Counting: Concepts, Algorithms, and Applications to Network Motifs and Graphlets. *Comput. Surveys* 54, 2 (2021), 1–36.
- [14] Siddhartha Sahu, Amine Mhedhbi, Semih Salihoglu, Jimmy Lin, and M. Tamer Özsu. 2020. The Ubiquity of Large Graphs and Surprising Challenges of Graph Processing: Extended Survey. *The VLDB Journal* 29 (2020), 595–618.
- [15] Volker Strassen. 1969. Gaussian Elimination is Not Optimal. *Numer. Math.* 13, 4 (1969), 354–356.
- [16] Jakub Tětek. 2022. Approximate Triangle Counting via Sampling and Fast Matrix Multiplication. In *49th International Colloquium on Automata, Languages, and Programming (ICALP 2022) (LIPIcs, Vol. 229)*. 107:1–107:20.